

6 Estimation

6.1 Introduction

The purpose of statistics is to make inferences about populations based on data from a small sample of that population. Populations are characterized by probability distributions which can be described by numerical parameters. The population mean μ and population variance σ^2 are parameters common to all probability distributions. In addition, some populations can be described by other natural parameters. One example is the population of all registered voters in Rhode Island. If we are interested in the yes/no question “Are you going to vote for Clinton?”, a natural parameter is p , the proportion of the population who plans on voting for Clinton.

The setup is exactly the same as with sampling. Suppose we are studying a population whose distribution is characterized by a parameter of interest which we will denote θ (this could be the population mean, population variance, proportion of voters supporting Clinton, or some other parameter). We will take n independent, identically distributed samples Y_1, \dots, Y_n from our population. An *estimator* is a function of our n samples which is designed to give us information about the population parameter. There are two types of estimators we will discuss:

1. A *point estimator* produces a single number which we think is close to the parameter of interest. We will learn several ways to quantitatively evaluate the “goodness” of a point estimator
2. An *interval estimate* produces an interval (often called a confidence interval) in which we believe our parameter of interest lies. We will learn how to construct confidence intervals which include the parameter of interest with a given probability.

6.2 Point Estimators

We start our discussion with point estimators. We have already met one point estimator, the sample mean \bar{Y} . The sample mean is an estimator since it is a function of our n samples. It is an estimator for the population mean. For another example, suppose we are polling n voters out of a population of registered voters and asking them if they are voting for Clinton. The parameter of interest is p , the proportion of registered voters who will vote for Clinton. Let Y be the number of voters in our sample who are voting for Clinton. Then the sample proportion $\hat{p} = Y/n$ is an estimator for the population proportion p . (The “hat” over the p indicates that it is an estimator for the parameter p).

Since we are applied mathematicians, we need quantitative tools to tell whether an estimator is any good. Let $\hat{\theta}$ be an estimator for the parameter θ . The first criterion we can use to evaluate an estimator is *bias*. We would like the expected value of our estimator to be the actual value of the parameter we are trying to estimate, i.e. $\mathbb{E}(\hat{\theta}) = \theta$. An estimator which has this property is called *unbiased*

Bias of an Estimator

Let $\hat{\theta}$ be an estimator for a parameter θ . The estimator $\hat{\theta}$ is *unbiased* if $\mathbb{E}(\hat{\theta}) = \theta$. Otherwise, the estimator $\hat{\theta}$ is *biased*. The *bias* of $\hat{\theta}$ is given by

$$\text{Bias}(\hat{\theta}) = \mathbb{E}(\hat{\theta}) - \theta$$

Let's look at the estimators we have seen so far. The sample mean \bar{Y} is an estimator for the population mean μ (here we abandon our “hat”-convention, and call this \bar{Y} rather than $\hat{\mu}$). We showed in the last section that $\mathbb{E}(\bar{Y}) = \mu$, thus the sample mean is an unbiased estimator for the population mean.

What about the sample proportion estimator we use in polling? Suppose we poll n voters in a population, and let Y be the number of voters in our sample who are voting for Clinton. As long as the sample is small enough (less than 1/20 of the population size), we can take Y to be a binomial random variable with parameters n and p . For our estimator $\hat{p} = Y/n$,

$$\mathbb{E}(\hat{p}) = \mathbb{E}\left(\frac{Y}{n}\right) = \frac{1}{n}\mathbb{E}(Y) = \frac{np}{n} = p$$

where we have used the expected value of a binomial random variable. Since $\mathbb{E}(\hat{p}) = p$, this estimator is unbiased as well. We will see an example of a biased estimator later.

A perhaps better measure of the “goodness” of an estimator is its *mean square error*, the average of the square distance of the estimator from the parameter of interest.

Mean Square Error

Let $\hat{\theta}$ be an estimator for a parameter θ . The *mean square error* of $\hat{\theta}$ is defined by

$$MSE(\hat{\theta}) = \mathbb{E}[(\hat{\theta} - \theta)^2]$$

If $\text{Bias}(\hat{\theta})$ is the bias of $\hat{\theta}$ and $\text{Var}(\hat{\theta})$ is the variance of $\hat{\theta}$, then

$$MSE(\hat{\theta}) = [\text{Bias}(\hat{\theta})]^2 + \text{Var}(\hat{\theta})$$

Note that the MSE can be divided into two components: variance and bias squared. Both of these are positive. We have discussed above how low bias is a good quality for an estimator. Low variance is a desired quality for an estimator as well since ideally we would like the distribution of the estimator to cluster tightly about the parameter of interest. In general, if we keep the sample size fixed, there is a tradeoff between bias and variance. For a given MSE, if we wish our estimator to have lower bias, then we must accept a higher variance

and vice versa.

To show the relationship between MSE, bias, and variance, we take the definition of MSE and add and subtract $\mathbb{E}(\hat{\theta})$ inside the parentheses.

$$\begin{aligned}
MSE(\hat{\theta}) &= \mathbb{E}[(\hat{\theta} - \theta)^2] \\
&= \mathbb{E}[(\hat{\theta} - \mathbb{E}(\hat{\theta})) + (\mathbb{E}(\hat{\theta}) - \theta)]^2 \\
&= \mathbb{E}[(\hat{\theta} - \mathbb{E}(\hat{\theta}))^2] + 2\mathbb{E}[(\hat{\theta} - \mathbb{E}(\hat{\theta}))\underbrace{(\mathbb{E}(\hat{\theta}) - \theta)}_{\text{constant}}] + \mathbb{E}[\underbrace{(\mathbb{E}(\hat{\theta}) - \theta)^2}_{\text{constant}}] \\
&= Var(\hat{\theta}) + 2(\mathbb{E}(\hat{\theta}) - \theta)\mathbb{E}[(\hat{\theta} - \mathbb{E}(\hat{\theta}))] + (\mathbb{E}(\hat{\theta}) - \theta)^2 \\
&= Var(\hat{\theta}) + 2(\mathbb{E}(\hat{\theta}) - \theta)\underbrace{(\mathbb{E}(\hat{\theta}) - \mathbb{E}(\hat{\theta}))}_{=0} + [Bias(\hat{\theta})]^2 \\
&= Var(\hat{\theta}) + [Bias(\hat{\theta})]^2
\end{aligned}$$

Let's look at the MSE of the two estimators we discussed above.

1. Sample mean $\bar{Y} = \frac{1}{n} \sum_{i=1}^n Y_i$

We showed that this estimator is unbiased, so the MSE is equal to the variance. We showed in a previous section that the variance of the sample mean is σ^2/n , where σ^2 is the population variance. Thus we have

$$MSE(\bar{Y}) = \frac{\sigma^2}{n}$$

Note that the MSE goes to 0 as $n \rightarrow \infty$, i.e. the error of our estimator decreases as our sample gets larger. This makes intuitive sense that a larger sample provides a better estimator for the population mean.

2. Sample proportion. $\hat{p} = \frac{Y}{n}$.

Recall that we are assuming that $Y \sim \text{Binomial}(n, p)$. We showed above that this estimator is also unbiased, thus once again the MSE is equal to the variance. Recalling that the variance of a binomial random variable is $np(1-p)$,

$$\begin{aligned}
MSE(\hat{p}) &= Var\left(\frac{Y}{n}\right) \\
&= \frac{1}{n^2} Var(Y) \\
&= \frac{np(1-p)}{n^2} \\
&= \frac{p(1-p)}{n}
\end{aligned}$$

Sometimes we are interested in studying the difference between two populations. Here are some examples of that.

1. Suppose we are interested in whether Brown University first-year students or seniors get more sleep. In this case, the parameter of interest is the *difference* in the mean amount of sleep between first-years and seniors. If μ_1 and σ_1^2 are the mean and variance of the amount of sleep of first-years and μ_2 and σ_2^2 are the same for seniors, then mathematically our parameter of interest is $\mu_1 - \mu_2$. If we take a sample of n_1 first-year students and n_2 seniors and compute the sample means \bar{Y}_1 and \bar{Y}_2 , then $\bar{Y}_1 - \bar{Y}_2$ is an estimator for $\mu_1 - \mu_2$. By linearity of expectation and our result for the expected value of the sample mean, the expected value of the sample mean is $\mu_1 - \mu_2$, so this estimator is unbiased. For the variance of this estimator, we use the formula for the variance of a sum (assuming the samples are independent), and recall that constants are squared when they are pulled out of the variance:

$$\begin{aligned} \text{Var}(\bar{Y}_1 - \bar{Y}_2) &= \text{Var}(\bar{Y}_1) + \text{Var}(-\bar{Y}_2) \\ &= \text{Var}(\bar{Y}_1) + (-1)^2 \text{Var}(\bar{Y}_2) \\ &= \text{Var}(\bar{Y}_1) + \text{Var}(\bar{Y}_2) \\ &= \frac{\sigma_1^2}{n_1} + \frac{\sigma_2^2}{n_2} \end{aligned}$$

2. Suppose we are interested in the preference for Clinton in rural versus urban voters in Pennsylvania. The parameter of interest here is the difference in the proportion of Clinton supporters between rural and urban areas. First, we have to define “rural” and “urban”. Suppose we take “urban” to mean living in a city with population of 250,000 or more (in Pennsylvania, this would only include Philadelphia and Pittsburgh). What are advantages or drawbacks to this definition? If p_1 is the proportion of Clinton supports in rural areas and p_2 is the proportion of Clinton supporters in urban areas, then the parameter of interest is $p_1 - p_2$. Suppose we sample n_1 voters from rural areas and n_2 voters from urban areas. Let Y_1 and Y_2 be the proportion of rural and urban voters (respectively) who support Clinton. Then

$$\hat{p}_1 - \hat{p}_2 = \frac{Y_1}{n_1} - \frac{Y_2}{n_2}$$

is an estimator for $p_1 - p_2$. By linearity of expectation and the result for a single population, the expected value of this estimator is $p_1 - p_2$, so this estimator is unbiased. What is the variance of this estimator? As above, using the formula for the variance of a sum of two independent random variables,

$$\begin{aligned} \text{Var}(\hat{p}_1 - \hat{p}_2) &= \text{Var}(\hat{p}_1) + \text{Var}(-\hat{p}_2) \\ &= \text{Var}(\hat{p}_1) + (-1)^2 \text{Var}(\hat{p}_2) \\ &= \text{Var}(\hat{p}_1) + \text{Var}(\hat{p}_2) \\ &= \frac{p_1(1 - p_1)}{n_1} + \frac{p_2(1 - p_2)}{n_2} \end{aligned}$$

We summarize these common estimators for population mean and proportion in the following table:

Parameter of Interest	Sample Size	Estimator	Expected Value	Variance	Standard Deviation
μ	n	\bar{Y}	μ	$\frac{\sigma^2}{n}$	$\frac{\sigma}{\sqrt{n}}$
p	n	$\hat{p} = \frac{Y}{n}$	p	$\frac{p(1-p)}{n}$	$\sqrt{\frac{p(1-p)}{n}}$
$\mu_1 - \mu_2$	n_1 and n_2	$\bar{Y}_1 - \bar{Y}_2$	$\mu_1 - \mu_2$	$\frac{\sigma_1^2}{n_1} + \frac{\sigma_2^2}{n_2}$	$\sqrt{\frac{\sigma_1^2}{n_1} + \frac{\sigma_2^2}{n_2}}$
$p_1 - p_2$	n_1 and n_2	$\hat{p}_1 - \hat{p}_2$	$p_1 - p_2$	$\frac{p_1(1-p_1)}{n_1} + \frac{p_2(1-p_2)}{n_2}$	$\sqrt{\frac{p_1(1-p_1)}{n_1} + \frac{p_2(1-p_2)}{n_2}}$

The standard deviation of an estimator is sometimes called the *standard error*. You may see that term in the scientific literature, but we will not use it in class. Note that the expected values and variances in the table above hold regardless of the distribution of the underlying population. In the case that the population is normal, the sample mean estimators \bar{Y} and $\bar{Y}_1 - \bar{Y}_2$ have a normal distribution, as discussed in the chapter on sampling distributions. However, for large sample sizes (approximately $n \geq 30$), the central limit theorem comes into play. Thus for large sample sizes, the sample mean estimators \bar{Y} and $\bar{Y}_1 - \bar{Y}_2$ are approximately normally distributed regardless of the distribution of the underlying population. For a binomial population, the population proportion estimators \hat{p} and $\hat{p}_1 - \hat{p}_2$ are also approximately normal for large sample sizes. How large a sample size do we need in this case? It depends on p . The farther p is from $1/2$, the larger sample size n we need for the binomial distribution to be approximately normal. In the homework, we showed that this is the case when $0 \leq 3\sqrt{pq/n} \leq 1$, where $q = 1 - p$. We then showed that an easier rule to check is that the binomial distribution is approximately normal when $n \geq 9p/q$ and $n \geq 9q/p$.

As another example, let's look at our estimators for population variance. Recall from the previous section that the sample variance is given by:

$$S^2 = \frac{1}{n-1} \sum_{i=1}^n (Y_i - \bar{Y})^2$$

where \bar{Y} is the sample mean. The factor of $(n-1)$ in the denominator seems peculiar. It seems more natural to divide by n and use the following estimator for the population variance:

$$S'^2 = \frac{1}{n} \sum_{i=1}^n (Y_i - \bar{Y})^2$$

We will show that while they are both estimators for the sample variance, S'^2 is biased whereas S^2 is unbiased. Since S^2 is unbiased, we call it the sample variance.

First we show S'^2 is biased by computing its expected value. This is done in several steps.

1. First we find a nice formula for $\sum_{i=1}^n (Y_i - \bar{Y})^2$.

$$\begin{aligned}
\sum_{i=1}^n (Y_i - \bar{Y})^2 &= \sum_{i=1}^n (Y_i^2 - 2Y_i\bar{Y} + \bar{Y}^2) \\
&= \sum_{i=1}^n Y_i^2 - 2\bar{Y} \sum_{i=1}^n Y_i + \sum_{i=1}^n \bar{Y}^2 \\
&= \sum_{i=1}^n Y_i^2 - 2n\bar{Y}^2 + n\bar{Y}^2 \\
&= \sum_{i=1}^n Y_i^2 - n\bar{Y}^2
\end{aligned}$$

2. Next we take the expected value of this. By linearity of expectation,

$$\begin{aligned}
\mathbb{E} \left[\sum_{i=1}^n (Y_i - \bar{Y})^2 \right] &= \mathbb{E} \left(\sum_{i=1}^n Y_i^2 \right) - n\mathbb{E}(\bar{Y}^2) \\
&= \sum_{i=1}^n \mathbb{E}(Y_i^2) - n\mathbb{E}(\bar{Y}^2)
\end{aligned}$$

3. Next we use the Magic Variance Formula “in reverse” to compute $\mathbb{E}(Y_i^2)$ and $\mathbb{E}(\bar{Y}^2)$.

$$\begin{aligned}
\text{Var}(Y_i) &= \mathbb{E}(Y_i^2) - [\mathbb{E}(Y_i)]^2 \\
\mathbb{E}(Y_i^2) &= \text{Var}(Y_i) + [\mathbb{E}(Y_i)]^2 \\
&= \sigma^2 + \mu^2
\end{aligned}$$

Similarly,

$$\begin{aligned}
\mathbb{E}(\bar{Y}^2) &= \text{Var}(\bar{Y}) + [\mathbb{E}(\bar{Y})]^2 \\
&= \frac{\sigma^2}{n} + \mu^2
\end{aligned}$$

4. We then plug these into the expression from step 3.

$$\begin{aligned}
\mathbb{E} \left[\sum_{i=1}^n (Y_i - \bar{Y})^2 \right] &= \sum_{i=1}^n (\sigma^2 + \mu^2) - n \left(\frac{\sigma^2}{n} + \mu^2 \right) \\
&= n\sigma^2 + n\mu^2 - \sigma^2 - n\mu^2 \\
&= (n-1)\sigma^2
\end{aligned}$$

5. Finally we use this result to compute the expected value of S'^2 .

$$\begin{aligned}\mathbb{E}(S'^2) &= \mathbb{E} \left[\frac{1}{n} \sum_{i=1}^n (Y_i - \bar{Y})^2 \right] \\ &= \frac{1}{n} \mathbb{E} \left[\sum_{i=1}^n (Y_i - \bar{Y})^2 \right] \\ &= \frac{1}{n} (n-1) \sigma^2 \\ &= \frac{n-1}{n} \sigma^2\end{aligned}$$

Since $\mathbb{E}(S'^2) \neq \sigma^2$, this estimator is biased. For large n , however, $(n-1)/n \approx 1$, so the bias is minimal. We can convert this to an unbiased estimator by multiplying by $n/(n-1)$. This is legitimate since n is a known constant (the sample size) and not one of the parameters we are trying to estimate:

$$\begin{aligned}\mathbb{E} \left(\frac{n}{n-1} S'^2 \right) &= \frac{n}{n-1} \mathbb{E}(S'^2) \\ &= \frac{n}{n-1} \frac{n-1}{n} \sigma^2 \\ &= \sigma^2\end{aligned}$$

So we have found an unbiased estimator for σ^2 . But we also have:

$$\begin{aligned}\frac{n}{n-1} S'^2 &= \frac{n}{n-1} \frac{1}{n} \sum_{i=1}^n (Y_i - \bar{Y})^2 \\ &= \frac{1}{n-1} \sum_{i=1}^n (Y_i - \bar{Y})^2 \\ &= S^2\end{aligned}$$

Thus S^2 is an unbiased estimator for the population variance σ^2 .

6.3 Interval Estimators

Once again, suppose we are studying a population whose distribution is characterized by a parameter of interest which we will denote θ . We will take n independent, identically distributed samples Y_1, \dots, Y_n from our population. An *interval estimator* or *confidence interval* uses the samples Y_1, \dots, Y_n to construct a confidence interval $[\hat{\theta}_L, \hat{\theta}_U]$. (The subscripts L and U denote the lower and upper endpoints of the interval, and the “hat” on the θ reminds us that this is an estimator). Note that the confidence interval is a random variable, and is a function of our n sample points. If we take a different sample, we will get a different confidence interval.

We would like the confidence interval to have the following properties:

1. It should contain the parameter of interest θ .
2. It should be relatively narrow (otherwise we haven't learned much).
3. We should be able to calculate the probability that our confidence interval will enclose our parameter of interest. This probability is called the *confidence coefficient*.

Suppose we have constructed a confidence interval $[\hat{\theta}_L, \hat{\theta}_U]$ for our parameter θ from our n samples Y_1, \dots, Y_n . Then the confidence coefficient is denoted $(1 - \alpha)$, and so we have¹:

$$\mathbb{P}(\hat{\theta}_L \leq \theta \leq \hat{\theta}_U) = 1 - \alpha$$

The confidence interval $[\hat{\theta}_L, \hat{\theta}_U]$ is sometimes called a *two-sided confidence interval*. We can also construct one-sided confidence intervals, although we will not do so in this course.


6.3.1 Confidence Intervals for Large Sample Sizes

Earlier we discussed the common unbiased estimators \bar{Y} (sample mean) and \hat{p} (sample proportion), as well as the equivalent estimators for the difference of two populations. For large sample sizes ($n \geq 30$ for the sample mean, and using our binomial rule for the sample proportion), the central limit theorem tells us that these estimators are all normally distributed with mean and standard deviation given in the table above. Thus, if we convert to the standard normal random variable, we can construct a confidence interval with desired confidence coefficient $(1 - \alpha)$ using the Z distribution. Let's see how we can do this.

Let θ be our parameter of interest (either $\mu, p, \mu_1 - \mu_2$, or $p_1 - p_2$), and let $\hat{\theta}$ be the appropriate unbiased estimator from the table above. In all cases, since the estimator is unbiased, the expected value $\mathbb{E}(\hat{\theta}) = \theta$. Let $\sigma_{\hat{\theta}}$ be the standard deviation of $\hat{\theta}$ (which we can look up in the table above). Then:

$$Z = \frac{\hat{\theta} - \theta}{\sigma_{\hat{\theta}}}$$

is (approximately) a standard normal random variable. Suppose we want a confidence interval for θ which has a confidence coefficient $(1 - \alpha)$. We will always use a two-sided, symmetric confidence interval (although you do not have to do this).



¹The reasons for the notation $(1 - \alpha)$ will be more evident when we discuss hypothesis testing; roughly speaking, α is the probability of a false positive result.

Converting to the standard normal random variable, this is equivalent to finding values $-z_{\alpha/2}$ and $z_{\alpha/2}$ such that $\mathbb{P}(-z_{\alpha/2} \leq Z \leq z_{\alpha/2}) = (1 - \alpha)$. This is illustrated in the picture above. To find $-z_{\alpha/2}$ and $z_{\alpha/2}$, we can use the Z table. For example, if $(1 - \alpha) = 0.95$, then $\alpha/2 = 0.025$. Consulting the Z table, we find that $-z_{\alpha/2} = -1.96$. By symmetry of the standard normal distribution about its mean of 0, $z_{\alpha/2} = 1.96$.

To find our confidence interval, we convert from the standard normal distribution back to the distribution of our estimator, which is $\text{Normal}(\theta, \sigma_{\hat{\theta}})$. Substituting for Z , we have

$$\mathbb{P}\left(-z_{\alpha/2} \leq \frac{\hat{\theta} - \theta}{\sigma_{\hat{\theta}}} \leq z_{\alpha/2}\right) = 1 - \alpha$$

Now we just manipulate this to get it into the form we want.

$$\begin{aligned}\mathbb{P}\left(-z_{\alpha/2}\sigma_{\hat{\theta}} \leq \hat{\theta} - \theta \leq z_{\alpha/2}\sigma_{\hat{\theta}}\right) &= 1 - \alpha \\ \mathbb{P}\left(-\hat{\theta} - z_{\alpha/2}\sigma_{\hat{\theta}} \leq -\theta \leq -\hat{\theta} + z_{\alpha/2}\sigma_{\hat{\theta}}\right) &= 1 - \alpha \\ \mathbb{P}\left(\hat{\theta} - z_{\alpha/2}\sigma_{\hat{\theta}} \leq \theta \leq \hat{\theta} + z_{\alpha/2}\sigma_{\hat{\theta}}\right) &= 1 - \alpha\end{aligned}$$

where to get the last line we multiplied by -1 and flipped all the inequalities. Thus the $(1 - \alpha)$ confidence interval for θ is:

$$[\hat{\theta}_L, \hat{\theta}_U] = [\hat{\theta} - z_{\alpha/2}\sigma_{\hat{\theta}}, \hat{\theta} + z_{\alpha/2}\sigma_{\hat{\theta}}]$$

We can also write this in terms of the population standard deviation σ and the sample size n by substituting $\sigma_{\hat{\theta}} = \sigma/\sqrt{n}$.

$$[\hat{\theta}_L, \hat{\theta}_U] = \left[\hat{\theta} - z_{\alpha/2}\frac{\sigma}{\sqrt{n}}, \hat{\theta} + z_{\alpha/2}\frac{\sigma}{\sqrt{n}}\right]$$

Let's do some examples.

Example. A ball bearing machine produces ball bearings whose diameters are normally distributed with mean μ mm and standard deviation $\sigma = 0.1$ mm. We would like to estimate the mean μ using a 90% confidence interval. To do this, we take a sample size of $n = 9$ ball bearings and use the sample mean \bar{Y} as our estimator for μ . Suppose we measure $\bar{Y} = 10.0$ mm. What is the 90% confidence interval for μ ?

Although this is not a large sample size, we can still find a confidence interval in this case because we are assuming that the population is normally distributed. Thus the sample mean \bar{Y} will always be normally distributed. Since we know the population standard deviation, we can convert to the standard normal Z distribution directly without appealing to approximations which require a large sample size. First compute the standard deviation of our estimator:

$$\sigma_{\bar{Y}} = \frac{\sigma}{\sqrt{n}} = \frac{0.1}{\sqrt{9}} = \frac{0.1}{3} = 0.0333$$

Since $1 - \alpha = 0.9$, $\alpha = 0.1$, thus $\alpha/2 = 0.05$. Looking at the Z table, we have to choose between a probability of 0.0495 ($z_{\alpha/2} = 1.65$) and a probability of 0.0505 ($z_{\alpha/2} = 1.64$). We will choose $z_{\alpha/2} = 1.65$ to allow for more “wiggle room”. Thus our 90% confidence interval for μ is given by:

$$\begin{aligned} [\bar{Y}_L, \bar{Y}_U] &= [\bar{Y} - z_{\alpha/2}\sigma_{\bar{Y}}, \bar{Y} + z_{\alpha/2}\sigma_{\bar{Y}}] \\ &= [\bar{Y} - (1.65)(0.0333), \bar{Y} + (1.65)(0.0333)] \\ &= [10.0 - (1.65)(0.0333), 10.0 + (1.65)(0.0333)] \\ &= [9.945, 10.055] \end{aligned}$$

Although we cannot say for certain that the true population parameter μ falls within this range, we can say that the probability that it does is 0.90. If we repeated this procedure, i.e. took 9 more samples, generated the sample mean, and constructed a 90% confidence interval, we would get a different confidence interval, but the probability that μ would fall in that confidence interval would still be 0.90.

Note that constructing this confidence interval required knowledge of the population standard deviation (which we then divided by n to get the standard deviation of the estimator). Often we do not know this parameter, but we wish to construct a confidence interval anyway. What do we do in that case? We use the sample standard deviation (found, for example, using the estimator S) in place of the population standard deviation. If this sounds to you like cheating, you are right! The key is that n is large. If the population is large, the sample standard deviation is close to the population standard deviation and there is very little loss of accuracy if we use S in place of σ in the formula for the confidence interval. We will justify this approximation mathematically later on, but for now just recall that the t distribution (which we use when the population standard deviation is unknown) approaches the standard normal Z distribution as the number of degrees of freedom increases.

Example. The shopping times of $n = 64$ randomly selected customers at a local supermarket were recorded. The sample mean and sample variance of the 64 shoppers were 33 minutes and 256 minutes, respectively. Find a 98% confidence interval for μ , the true average shopping time per customer.

In this case, the parameter of interest is μ . Our estimator is \bar{Y} , the sample mean. In this experiment, we sampled $n = 64$ customers and found a sample mean $\bar{Y} = 33$ and a sample variance $S^2 = 256$. Since the sample is large (more than 30 customers), by the central limit theorem, we can assume that \bar{Y} is normally distributed. We do not know the population variance, so we will use the sample variance in place of the population variance. Since the population is large, this is a reasonable assumption. Thus we can use the formula above for the confidence interval, where for the standard deviation of the estimator we use:

$$\sigma_{\bar{Y}} = \frac{S}{\sqrt{n}} = \frac{\sqrt{256}}{\sqrt{64}} = \frac{16}{8} = 2$$

To find $z_{\alpha/2}$, we look at our Z table. Since $(1 - \alpha) = 0.98$, $\alpha = 0.02$, thus $\alpha/2 = 0.01$. Looking at our Z table, the closest probability we have to 0.01 is 0.099 (0.0102 is the next

closest, but since it is farther away, we will use 0.099), and the value of Z corresponding to that gives us $z_{\alpha/2} = 2.33$. We now have everything we need to construct our 98% confidence interval.

$$\begin{aligned} [\bar{Y}_L, \bar{Y}_U] &= [\bar{Y} - z_{\alpha/2}\sigma_{\bar{Y}}, \bar{Y} + z_{\alpha/2}\sigma_{\bar{Y}}] \\ &= [33 - (2.33)(2), \bar{Y} + (2.33)(2)] \\ &= [28.34, 37.66] \end{aligned}$$

Here is another example, this time dealing with the difference between sample proportions.

Example. Two brands of lightbulbs, denoted brand A and brand B, are guaranteed to last for at least 1 year. In a random sample of 50 lightbulbs of brand A, 12 were found to fail before the 1 year period ended. In an independent random sample of 60 lightbulbs of brand B, 12 were also found to fail before the 1 year period ended. Give a 98% confidence interval for the difference $p_1 - p_2$ between the proportion of failures of the two brands during the 1 year period.

The parameter of interest here is $p_1 - p_2$, and we use as our estimator $\hat{p}_1 - \hat{p}_2$. Evaluating our estimators, we get $\hat{p}_1 = 12/50 = 0.24$ and $\hat{p}_2 = 12/60 = 0.20$, thus $\hat{p}_1 - \hat{p}_2 = 0.24 - 0.20 = 0.04$. To construct our confidence interval, we need the standard deviation of our estimator. From the table above (or by deriving it from the variance of the binomial distribution), we have

$$\sigma_{\hat{p}_1 - \hat{p}_2} = \sqrt{\frac{p_1(1 - p_1)}{n_1} + \frac{p_2(1 - p_2)}{n_2}}$$

But we have a problem. The standard deviation involves the true parameter p , and that is what we are trying to estimate! No worries, we will just do what we did before. Since the population is large, we will use the estimator \hat{p} in place of p in the expression for the standard deviation.

$$\begin{aligned} \sigma_{\hat{p}_1 - \hat{p}_2} &\approx \sqrt{\frac{\hat{p}_1(1 - \hat{p}_1)}{n_1} + \frac{\hat{p}_2(1 - \hat{p}_2)}{n_2}} \\ &= \sqrt{\frac{(0.24)(0.76)}{50} + \frac{(0.20)(0.80)}{60}} \\ &= 0.0795 \end{aligned}$$

We found $z_{\alpha/2}$ for a 98% confidence interval in the previous example, thus we have $z_{\alpha/2} = 2.33$. Using the formula, our 98% confidence interval is:

$$\begin{aligned} [(\hat{p}_1 - \hat{p}_2) - z_{\alpha/2}\sigma_{\hat{p}_1 - \hat{p}_2}, (\hat{p}_1 - \hat{p}_2) + z_{\alpha/2}\sigma_{\hat{p}_1 - \hat{p}_2}] &= [0.04 - (2.33)(0.0795), 0.04 + (2.33)(0.0795)] \\ &= [0.04 - 0.185, 0.04 + 0.185] \\ &= [-0.145, 0.225] \end{aligned}$$

Note that our 98% confidence overlaps 0. If the true parameter $p_1 - p_2$ were 0, then there would be no difference in the 1-year failure rates of the two brands of lightbulbs! Thus we cannot exclude that possibility with 98% confidence.

6.3.2 Experimental Design: Selecting the Sample Size

If you are designing an experiment, one of the parameters you must choose is the sample size. As an example, suppose you are polling registered voters in Rhode Island and asking them if they are voting for Clinton. You use the estimator \hat{p} for p , the true proportion of Clinton voters. As the number n of people polled (the size of your sample) increases, the variance of your estimator \hat{p} decreases, and there is a higher probability that your estimator will be close to the true proportion. In terms of confidence intervals, if you desire 95% confidence that your interval contains the true parameter p , as you increase the sample size, the confidence interval becomes narrower. There is a tradeoff, however, as the sample size gets larger. Collecting a larger sample is more expensive and consumes more resources. Is there any way mathematically to determine ahead of time the sample size we need to attain a desired level of accuracy for our estimator?

Suppose we want to estimate the average diameter μ of a ball bearing produced by a ball bearing machine, and we wish the probability that our estimate is within 0.02 mm of the true mean to be 0.95. Assuming the diameters of the ball bearings are normally distributed (or taking a large enough sample that the central limit theorem applies), by the 68-95-99.7 rule we want the true value μ to be within 2 standard deviations of our estimator \bar{Y} (This is the standard deviation of the *estimator*, not the population standard deviation). Thus we want:

$$2\sigma_{\bar{Y}} = 2\frac{\sigma}{\sqrt{n}} = 0.02$$

We can solve the above equation for n as long as we know σ .

$$n = \left(\frac{2\sigma}{0.02} \right)^2$$

What do we do if the population standard deviation σ is not known, as is often the case. One possibility is to use the sample standard deviation S from a previously taken sample in place of the population standard deviation σ . As long as we take a large enough sample, this is a reasonable estimate. Another alternative is to use the range (largest value minus smallest value) of a previous sample. If we assume that the population is roughly normally distributed, then we know from the 68-95-99.7 rule that 95% of the population will lie within two standard deviations of the mean. If our sample is large enough, then the sample range should be approximately 4σ , since we are assuming the range represents two standard deviations both above and below the mean. Thus we can approximate σ as one-fourth of the range.

In our ball bearing example, suppose in a previous sample of 50 ball bearings we had a range of 0.4 mm. Dividing this by 4, we can estimate $\sigma = 0.4/4 = 0.1$. Plugging this info into our formula above, we get:

$$n = \left(\frac{2 \cdot 0.1}{0.02} \right)^2 = 10^2 = 100$$